EXPLORATORY DATA ANALYSIS CSC 3220 - Real Estate Data

Robert Bingham
Phonethep Nakhonekhong
Eli Parker
John Taylor
Johnathan Rich

2022-11-08

Introduction

Our group, 5 little minds, have decided to look at real estate market data from 2012 to 2021 in order to predict future trends in the next 5 years. This data includes median list price for a given neighborhood in each state; median prices based upon housing unit categories, such as apartments, single-family housing, and condos; and year-to-year increases in sale prices of each unit.

Format the data

For the purposes of this assignment, were are turning off all warnings and centering each graph.

```
knitr::opts_chunk$set(echo = TRUE, warning = FALSE, message = FALSE, fig.align =
↪  'center')
```

Import the necessary libraries

```
library("ggplot2")
library("DT")
library("pander")
library("corrplot")
library("zoo")
library("reshape")
library("scales")
library("tidyverse")
```

Import the Data

We decided to use the data from this url from Kaggle for our dataset:

https://www.kaggle.com/datasets/thuynyle/redfin-housing-market-data?select=state_market_tracker.tsv000

```
state_market.df <- read.table("../data/state_market_tracker.tsv000", sep = '\t', header =
↪   TRUE)
```

Data Manipulation

Here, we have made R recognize the variables in the dataset that pertain to specific days, (i.e, 9/21/2022) as actual dates using the built-in as.Date function. We have also divided the median sale price and list price of homes in each neighborhood by 1000 in order to make the data more readable in subsequent graphs.

```
state_market.df$period_begin <- as.Date(state_market.df$period_begin)
state_market.df$period_end <- as.Date(state_market.df$period_end)
state_market.df$median_sale_price <- state_market.df$median_sale_price / 1000
state_market.df$median_list_price <- state_market.df$median_list_price / 1000
state_market.df$property_type[state_market.df$property_type == "Multi-Family (2-4 Unit)"]
↪   <- "Multi-Family"
```

Out of the 48 variables in the dataset, these are the ones we will consider (the important variables) in our subsequent exploratory analysis:

- median sale price - **Median sale price the housing units of each of the 27,054 neighborhoods**
- median list price - **Median list price the housing units of each of the 27,054 neighborhoods**
- median sale price year-on-year - **Year-to-year percentage increase in median sale price of each neighborhood**
- median list price year-on-year - **Year-to-year percentage increase in median list price of each neighborbood**
- median sale ppsf - **Median sale price per-square-foot of each neighborhood**
- median list ppsf - **Median sale price per-square-foot of each neighborhood**
- homes sold year-old-year - **Increase in number of homes sold in each neighborhood year-on-year**
- pending sales year-on-year - **Increase in number of pending sales in each neighborhood year-on-year**
- new listings year-on-year - **Increase in number of new listings in each neighborhood year-on-year**
- inventory year-on-year - **Increase in number of all listings in each neighborhood year-on-year**

Interpreting the Data

Here, we created a boxplot graph for the median list price of the homes in each neighborhood in the dataset in the year 2021, with each boxplot representing each state. Hawaii and and Connecticut were had the highest number of outliers. Since there were more outliers above the boxplot than below it in both cases, we attributed these anomalies to the high cost of living in both states, Since Hawaii is a vacation destination and Connecticut is a New England state with a close proximity to New York City. We can also assume that the states that have a lower overall median sale price also have more consistent prices since boxplots such as Oklahoma and Ohio have a much lower variance and standard deviation than states with higher medians.

```
split_by_year <- split(state_market.df, format(state_market.df$period_begin, "%Y"))
ggplot(split_by_year[[length(split_by_year)]], aes(x = state_code, y = median_sale_price,
↪   color=median_sale_price)) +
  geom_boxplot(colour="#003366", fill="#66FFFF", alpha=5) +
  geom_point() +
```

```
xlab("2021") +
ylab("Median Sales Price (in thousands)") +
ggtitle("Distribution of Median Sale Price per State") +
theme(plot.title = element_text(hjust = 0.5))
```


Distribution of Median Sale Price per State

We created 2 histograms. One is for the median list price of housing units in each of the 27,054 neighborhoods from 2012 to 2021. The other is the median sale price of over the same length of time. We noticed that both graphs are skewed to the right, so we can assume that there are more homes sold and listed above the overall median home price than below it in the U.S. We also noticed that there is much less variance in the sale price than in the list price, with the right skew in the list price being much less pronounced in the list price than the sale price. Therefore, we can assume that if the asking/list price is higher than the median, the seller receives less in the final sale price in most cases.

```
mean_years <- data.frame(state_market.df$period_begin,
↪   state_market.df$median_sale_price_yoy, state_market.df$state)

new_df <- subset.data.frame(state_market.df, select = c(state_code, median_list_price,
↪   median_sale_price),  drop = FALSE)

print(
    ggplot(new_df, aes_string(x=new_df$median_list_price))
    + geom_histogram(
      colour="darkorchid4", fill="darkorchid1", position="identity", bins=30, alpha=0.2
    )
    + ggtitle(paste("Frequency Distribution of median list price", sep=""))
    + theme(plot.title=element_text(hjust = 0.5))
    + xlab("Median List Price of Each Neighboorhood (in thousands)")
    + ylab("Frequency of Neighborhoods")
    + ylim(0, 12500))
```

Frequency Distribution of median list price

```
print(
    ggplot(new_df, aes_string(x=new_df$median_sale_price))
    + geom_histogram(
        colour="red", fill="firebrick1", position="identity", bins=30, alpha=0.2
    )
    + ggtitle(paste("Frequency Distribution of median sale price", sep=""))
    + theme(plot.title=element_text(hjust = 0.5))
    + xlab("Median Sale Price of Each Neighboorhood (in thousands)")
    + ylab("Frequency of Neighborhoods")
    + ylim(0, 12500))
```

## Frequency Distribution of median sale price

We created a correlation matrix with our important variables. We expected a strong positive correlation between median sale price and median list price, and an even stronger correlation between median sale price and median price-per-square-foot, but a feature that surprised us was that there was a modest negative correlation between median list price year-on-year and the pending sales year-on-year, which directly contradicts the basic supply-and-demand principle. However, we believe we can attribute this to people being less likely to buy a house if the house is more expensive.

```
cor.df <- subset(state_market.df, state_code == 'CA')

cor.df <- subset.data.frame(cor.df,  select = c(median_sale_price, median_list_price,
↪  median_sale_price_yoy, median_list_price_yoy, median_ppsf, median_list_ppsf,
↪  homes_sold_yoy, pending_sales_yoy, new_listings_yoy,
↪  inventory_yoy,sold_above_list_yoy,avg_sale_to_list_yoy))

cor.table <- cor(cor.df, use="pairwise.complete.obs")

corrplot(cor.table, type="upper")
```

```r
variables <- colnames(cor.df)
numeric_vars <- variables[c(-0,-11)]
pander(numeric_vars)
```

median_sale_price, *median_list_price*, *median_sale_price_yoy*, *median_list_price_yoy*, *median_ppsf*, *median_list_ppsf*, *homes_sold_yoy*, *pending_sales_yoy*, *new_listings_yoy*, *inventory_yoy* and *avg_sale_to_list_yoy*

```
model2 <- lm(
  pending_sales_yoy~inventory_yoy+new_listings_yoy, cor.df[numeric_vars]
)

summary(model2)
```

```
##
## Call:
## lm(formula = pending_sales_yoy ~ inventory_yoy + new_listings_yoy,
##     data = cor.df[numeric_vars])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.37185 -0.05858 -0.00910  0.04636  0.52399
##
## Coefficients:
##                   Estimate Std. Error t value Pr(>|t|)
## (Intercept)       0.024542   0.004106   5.977 3.91e-09 ***
## inventory_yoy    -0.358371   0.015643 -22.909  < 2e-16 ***
## new_listings_yoy  0.978292   0.025793  37.929  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09838 on 597 degrees of freedom
## Multiple R-squared:  0.7321, Adjusted R-squared:  0.7312
## F-statistic: 815.7 on 2 and 597 DF,  p-value: < 2.2e-16
```

```
par(mfrow=c(2,2))
plot(model2)
```

Residuals vs Fitted

Residuals

0.2

−0.4

21420
21437
25869

0.0   0.5   1.0

Fitted values

Normal Q−Q

Standardized residuals

4

0

−4

21437
21420

−3 −2 −1   0   1   2   3

Theoretical Quantiles

Scale−Location

√|Standardized residuals|

1.5

0.0

21428
21437
4464

0.0   0.5   1.0

Fitted values

Residuals vs Leverage

Standardized residuals

4

0

−4

4464
9075
17243
Cook's distance

1

0.5

0.5

0.00 0.02 0.04 0.06 0.08

Leverage

9

```
sample_size <- nrow(cor.df)
mod <- lm(pending_sales_yoy~new_listings_yoy+inventory_yoy, cor.df[numeric_vars])
cooksd <- cooks.distance(mod)
influential <- as.numeric(names(cooksd)[(cooksd > (4/sample_size))])
df2 <- cor.df[-influential, ]
model4 <- lm(pending_sales_yoy~new_listings_yoy+inventory_yoy, df2[numeric_vars])
par(mfrow=c(2,2))

summary(model4)
```

```
##
## Call:
## lm(formula = pending_sales_yoy ~ new_listings_yoy + inventory_yoy,
##     data = df2[numeric_vars])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.37160 -0.05855 -0.00902  0.04626  0.52422
##
## Coefficients:
##                   Estimate Std. Error t value Pr(>|t|)
## (Intercept)       0.024333   0.004107    5.925 5.29e-09 ***
## new_listings_yoy  0.978207   0.025779   37.946  < 2e-16 ***
## inventory_yoy    -0.358359   0.015635  -22.921  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09833 on 596 degrees of freedom
## Multiple R-squared:  0.7326, Adjusted R-squared:  0.7317
## F-statistic: 816.5 on 2 and 596 DF,  p-value: < 2.2e-16
```
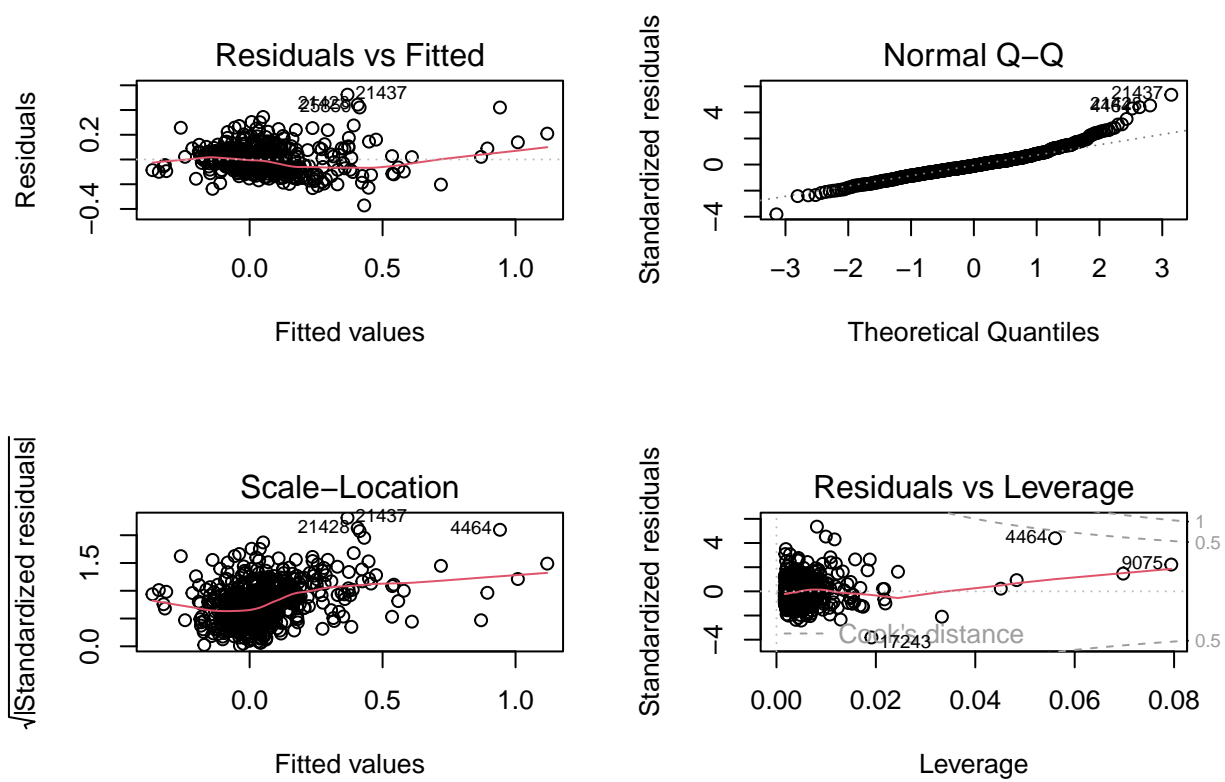
```
plot(model4)
```

Residuals vs Fitted

Residuals

0.2

−0.4

21437
21428
25839

0.0   0.5   1.0

Fitted values

Normal Q−Q

Standardized residuals

4

0

−4

21437
21428

−3 −2 −1   0   1   2   3

Theoretical Quantiles

Scale−Location

√|Standardized residuals|

1.5

0.0

21437
21428
25839

4464

0.0   0.5   1.0

Fitted values

Residuals vs Leverage

Standardized residuals

4

0

−4

4464

9075

17243

Cook's distance

1
0.5

0.5

0.00 0.02 0.04 0.06 0.08

Leverage

We can create a bin plot to demonstrate that home sale prices tend to aggregate below a million, between 100 thousand to 300 thousand. We have fitted a best fit line to show that there is a positive increase in median sales price over time, however, with a r-squared value of 0.125, only 12.5% of this increase can be explained by time. This is understandable, since the value of a home includes many factors such as: location, land, size, time built, etc. We can see that both the townhouse and multi-family units have extreme outliers in the median sales price.

```
lm_eqn <- function(df){
    m <- lm(median_sale_price ~ as.yearmon(period_begin), df);
    eq <- substitute(italic(median_sale_price) == b %.% italic(period_begin) +
↪    a*","~~italic(r)^2~"="~r2,
        list(a = format(unname(coef(m)[1]), digits = 2),
             b = format(unname(coef(m)[2]), digits = 2),
            r2 = format(summary(m)$r.squared, digits = 3)))
    as.character(as.expression(eq));
}

ggplot(state_market.df, aes(x = as.yearmon(period_begin), y = median_sale_price)) +
  geom_bin_2d(binwidth = c(1/12, 10)) +
  xlab("Time") +
  ylab("Median Sales Price") +
  ggtitle("Median Sales Price vs Time") +
  theme(plot.title = element_text(hjust = 0.5)) +
  scale_x_yearmon(n = 10) +
  geom_smooth(method = "lm", se = FALSE, color = "red") +
  annotate("text", x = as.numeric(as.yearmon("2014-06-01")), y = 3000, parse = TRUE,
↪    label = lm_eqn(state_market.df))
```

Median Sales Price vs Time

$median\_sale\_price = 17 \cdot period\_begin + -34977, \; r^2 = 0.125$

Below, we have created 2 boxplots for the types of housing units. One for the median sale price, the other for the median list price.

```
ggplot(state_market.df, aes(x = property_type, y = median_sale_price)) +
  geom_boxplot(colour="#003366", fill="#66FFFF", alpha=1/2) +
  xlab("Property Type") +
  ylab("Median Sales Price (in thousands)") +
  ggtitle("Property Type vs Median Sales Price") +
  theme(plot.title = element_text(hjust = 0.5))
```

Both townhouse and multi-family have extreme outliers in the median sales price.


```
## For our outlier tests, we will test if anything is above the 99th quantile. We can see
## that for the outliers for multi-family housing the state Hawaii has the most outliers. This
## gives us a good explanation in the reason for the outliers, since Hawaii is a small state
## in the middle of the Pacific serverely limiting the supply of land.

##
## California    Columbia       Hawaii New Mexico
##         11          11           30          3

## The same thing is repeated with townhouses, but Connecticut has almost as many as Hawaii
```

```
## now. There is a possibility that the reasoning behind this is the same as Hawaii with the
## limited supply of land, but it also could be that mostly everywhere in Connecticut is close
## to a town or city that sports many amenities.


##
##    Columbia Connecticut      Hawaii        Maine
##          9           19          20            2
```

```
ggplot(state_market.df, aes(x = property_type, y = median_list_price)) +
  geom_boxplot(colour="#003366", fill="#66FFFF", alpha=1/2) +
  xlab("Property Type") +
  ylab("Median Sales Price (in thousands)") +
  ggtitle("Property Type vs Median List Price") +
  theme(plot.title = element_text(hjust = 0.5))
```

Below, we can see that the extreme outliers are not as pronounced with the median list price.


```
## California leads the US in most outliers for all residential house prices while Columbia
## and Hawaii are not far behind. Hawaii has the same reasoning for the listing price as what
## was explained above. California's outliers can be explained by the housing markets
## predictions that a huge demand for housing will always be present in the state. Columbia has
## a low supply compared to the high demand in the housing market.
```

```
##
## California   Columbia     Hawaii      Maine
##         22        19         15          1
```

```
## New York leads the US in most outliers for condo/co-op house prices. New York will
## obviously lead in most outliers, since most people only live in condo's or cooperative
## housing in the state and the housing markets prediction of ever increasing demands of housing
## in places like New York City.
```

```
##
## California   New York       Utah
##          4        49          2
```

```
## California and Columbia leads the US in most outliers for single residential house prices.
## California is the same as described above. Columbia has a low supply compared to the high demand
## in the housing market.
```

```
##
## California   Columbia     Hawaii      Maine
##         19        23         14          1
```

Conclusion

We are going to use linear regression to predict the median sale price of different housing units in 5 years. Since our R^2 is low in the binplot equation, we know that only 12.5% of the independent variable (median sales price) can be explained through time. We can see that this is a small, but reliable relationship. For us to predict future house prices accurately, we will need to account for other factors that are outside the scope of this dataset; however, we may be able to predict a lower or upper bound or a relative area that future house prices might fall into. Either way, we can confidently confirm that the future housing market will be costlier that the current one.

Sample Data

```
sample.df <- subset.data.frame(state_market.df, select = c(period_begin, state_code,
→  property_type, median_sale_price, median_list_price, median_sale_price_yoy,
→  median_list_price_yoy, median_ppsf, median_list_ppsf, homes_sold_yoy,
→  pending_sales_yoy, new_listings_yoy, inventory_yoy),  drop = FALSE)

pander(head(sample.df))
```

Table 1: Table continues below

period_begin	state_code	property_type	median_sale_price
2019-10-01	OK	Multi-Family	162.2
2021-07-01	VT	All Residential	317.9
2016-08-01	NH	Condo/Co-op	200.1
2013-04-01	MS	All Residential	129.5
2019-12-01	MO	Condo/Co-op	152
2019-07-01	NM	All Residential	385.5

Table 2: Table continues below

median_list_price	median_sale_price_yoy	median_list_price_yoy
185	0.06697	0.005355
322.7	0.2048	0.07324
260.2	0.08944	0.09914
144.2	0.07135	0.05368
170.1	0.068	0.05685
390.8	0.08515	-0.05731

Table 3: Table continues below

median_ppsf	median_list_ppsf	homes_sold_yoy	pending_sales_yoy
77	107	-0.02326	0.175
177	183	-0.07414	0.1391
155	145	0.04545	0.0101
69	79	0.2043	0.1449
135	122	0.1264	NA
218	224	0.1007	-0.01429

new_listings_yoy	inventory_yoy
-0.1875	-0.2957
-0.0911	-0.27
-0.08157	-0.218
0.05665	0.2042
-0.09444	-0.2632
-0.1557	0.02579